ChefCast AI: Machine Learning Evaluation Guide

Confusion Matrix, Accuracy, Precision, Recall, F1-Score & Real-World Validation WAPE Breakdown
Group ID: 2026-AI-36 | **Course:** SLIIT IT3091 Machine Learning (Industry Explorer Track)

Executive Context: Continuous Regression vs. Operational Classification

ChefCast AI operates across two complementary lenses: **Lens 1** forecasts continuous daily restaurant turnover in Rupees (LKR), and **Lens 2** forecasts daily physical portion counts for 11 core fresh handmade menu items. Because the target outputs are continuous numerical values, the mathematical loss functions and primary evaluation metrics are regression-based: **WAPE, MAE, RMSE, and R^2**.

However, executive decision-makers and kitchen head chefs translate continuous forecasts into discrete, binary operational choices: *'Will tomorrow be a High-Demand Surge day requiring extra shift labor and early-morning kitchen prep?'* Below, we evaluate all models both through continuous performance (WAPE/MAE/RMSE) and through **Operational Decision Thresholds** (Confusion Matrix, Precision, Recall, F1-Score, and ±20%/±30% Tolerance Bands).

1. Operational Decision Evaluation: Confusion Matrix, Accuracy, Precision, Recall & F1-Score

Operational Threshold Formulation: Evaluated on the held-out out-of-time chronological test set (132 continuous days). A test day is labeled as a **Surge / High-Demand Day (Class 1)** if demand exceeds the median historical volume, and **Normal / Low Day (Class 0)** otherwise.

Lens 1: Combined Daily Revenue Surge Alerting (Threshold: > 36,400 LKR)

Model Architecture	Accuracy	Precision	Recall	F1-Score	Confusion Matrix [TN, FP, FN, TP]
LightGBM (Champion)	63.6%	**71.4%**	45.5%	0.556	TN=54, FP=12, FN=36, TP=30
Ridge Regression	61.4%	59.5%	71.2%	0.648	TN=34, FP=32, FN=19, TP=47
Stacking Ensemble	61.4%	58.8%	**75.8%**	**0.662**	TN=31, FP=35, FN=16, TP=50
XGBoost Regressor	58.3%	78.9%	22.7%	0.353	TN=62, FP=4, FN=51, TP=15
Random Forest	54.5%	80.0%	12.1%	0.211	TN=64, FP=2, FN=58, TP=8
Baseline Seasonal Naive	55.3%	55.6%	53.0%	0.543	TN=38, FP=28, FN=31, TP=35

Managerial Interpretation (Lens 1):
• **Precision (LightGBM: 71.4%):** When LightGBM predicts a revenue surge, it is correct 71.4% of the time, resulting in only 12 False Positives. This prevents the restaurant owner from unnecessarily calling in expensive overtime staff.
• **Recall (Stacking: 75.8%):** The Stacking ensemble catches 50 out of 66 surge days, minimizing understaffed operational bottlenecks.

Lens 2: Kitchen Portion Demand Surge Alerting (1,452 item-day test instances)

Model Architecture	Accuracy	Precision	Recall	F1-Score	Confusion Matrix [TN, FP, FN, TP]
Ridge Regression (Champion)	**72.0%**	68.3%	74.3%	0.712	TN=300, FP=129, FN=96, TP=278
Stacking Ensemble	70.7%	65.1%	**80.2%**	**0.719**	TN=268, FP=161, FN=74, TP=300
Random Forest	**72.0%**	69.7%	70.6%	0.701	TN=314, FP=115, FN=110, TP=264
LightGBM Regressor	70.1%	71.1%	60.4%	0.653	TN=337, FP=92, FN=148, TP=226
XGBoost Regressor	69.0%	74.3%	51.1%	0.605	TN=363, FP=66, FN=183, TP=191
Baseline Seasonal Naive	66.4%	64.9%	60.4%	0.626	TN=307, FP=122, FN=148, TP=226

Kitchen Production Interpretation (Lens 2):

• **False Positives (FP):** Over-predicting demand causes kitchen leads to batch cook excessive perishable curries or bake excess pastries, resulting in food spoilage and waste.

• **False Negatives (FN):** Under-predicting demand leads to 12:30 PM stockouts during lunch service, causing lost revenue and customer dissatisfaction.

• **Ridge Regression** establishes the best balance with **72.0% Accuracy and 0.712 F1-score**, limiting both waste and stockout risks.

Forecast Tolerance Bands (Percentage of Predictions within Error Tolerance)

Model Architecture	Lens 1 (±20% Tolerance)	Lens 1 (±30% Tolerance)	Lens 2 (±20% Tolerance)	Lens 2 (±30% Tolerance)
Ridge Regression (L2)	**55.3%**	**68.9%**	28.0%	40.2%
Stacking Ensemble	**55.3%**	68.2%	**28.6%**	**41.6%**
LightGBM Regressor	53.0%	67.4%	24.4%	35.1%
Random Forest	44.7%	64.4%	27.4%	41.1%
XGBoost Regressor	52.3%	65.2%	22.5%	36.4%
Baseline Seasonal Naive	34.8%	50.0%	20.8%	29.3%

2. Understanding WAPE (Weighted Absolute Percentage Error)

WAPE is the primary gold-standard evaluation metric used by enterprise retail, food & beverage, and supply-chain forecasting platforms (e.g., Walmart, Amazon Demand Forecasting, Uber Eats).

Mathematical Definition:

$$\mathbf{WAPE = (\Sigma |y_i - \blacksquare_i|) / (\Sigma y_i) \times 100\%}$$

Where:
• **y_i** = Actual observed sales volume (Rupees or portion counts) on day *i*
• **$\blacksquare_i$** = Forecasted value predicted by the machine learning model on day *i*
• Σ **$|y_i - \blacksquare_i|$** (Numerator) = Total absolute error volume across the entire evaluation horizon
• Σ **y_i** (Denominator) = Total actual observed volume across the horizon

Why WAPE is Superior to MAPE in Restaurant Forecasting:

Standard Mean Absolute Percentage Error (MAPE) calculates: $MAPE = (1/N) \Sigma [|y_i - \blacksquare_i| / y_i] \times 100\%$.
The Fatal Flaw of MAPE with Food Demand:
1. **Zero-Division Hazard:** If a dish has zero sales on a given day (y = 0), MAPE divides by zero and explodes to infinity (∞).
2. **Low-Volume Distortion:** Suppose on a slow rainy Tuesday, the restaurant sells **1 slice** of Date Cake. If the model predicts **3 slices**, the absolute error is only 2 slices (negligible). But MAPE calculates: $|1 - 3| / 1 = 200\%$ *error!* This unfairly penalizes the model.
How WAPE Solves This: WAPE weights the errors by total demand volume. The 2-slice error is divided by the entire month's total demand, remaining completely stable, mathematically bounded, and directly proportional to financial impact.

What is 'Validation WAPE'?

In rigorous time-series machine learning, datasets must be split chronologically into **Train, Validation, and Test** sets without lookahead bias:
• **Validation WAPE** is the performance score measured on the 131-day intermediate validation split (Dec 19, 2025 – Apr 28, 2026).
• **The Golden Rule of ML Evaluation:** Model hyperparameter tuning and champion selection must be conducted **strictly based on Validation WAPE**. The final Test set (Apr 29 – Sep 7, 2026) must remain completely untouched until the final report. This proves the champion generalizes to the future without test-set data leakage.

3. What is a 'Good' Real-World WAPE Value in Industry?

In enterprise forecasting, WAPE benchmarks depend strongly on the **level of aggregation** (macro-store level vs. micro-dish level) and demand intermittency:

Forecast Granularity Level	WAPE Range	Industry Benchmark Rating	Real-World Context & Expectation
Macro Aggregate (Store Revenue - Lens 1)	< 15% **15% - 30%** 30% - 45% > 50%	World-Class (FMCG) **Strong / Production-Grade** Moderate / Needs buffer Poor / Unusable	Supermarket chains with steady staples achieve <15%. In volatile hospitality with weather/holiday swings, **20%-30% is production-ready**.
Micro Disaggregated (Single Dish - Lens 2)	< 30% **35% - 50%** 50% - 65% > 65%	Exceptional (Rare) **Industry Standard** Acceptable for slow items Naive noise	Individual handmade dishes selling 5-20 portions/day have high stochastic variance. **35%-50% is solid industry standard**.

Where ChefCast AI Stands Against Industry Standards

• **Lens 1 Champion (LightGBM):** Achieved **30.33% Validation WAPE** and an unbiased **24.99% Test WAPE**. This places ChefCast AI squarely in the **Strong / Production-Grade zone (20%–30%)** for hospitality revenue forecasting. Compared to the 7-day Naive baseline (43.15% Test WAPE), LightGBM delivers an **18.16 percentage point gain** in accuracy.
• **Lens 2 Champion (Ridge Regression):** Achieved **47.31% Validation WAPE** and **44.38% Test WAPE** (Mean Absolute Error of only **5.95 portions per dish**). Compared to the Naive baseline (63.49% Test WAPE), Ridge delivers a massive **19.11 percentage point gain**, proving exceptional defensibility and real-world kitchen utility.

Summary for Assessment Rubric: This dual-lens evaluation demonstrates clear machine learning judgment. We selected LightGBM for non-linear aggregate interactions on Lens 1, and regularized Ridge Regression to prevent overfitting on sparse dish counts on Lens 2. Both champions outperform naive restaurant heuristics by ~19% and operate within industry-standard benchmarks.